

```

struct device {
    struct kobject kobj;
    struct device *parent;

    struct device private *p;

    const char *init_name; /* initial name of the device */
    const struct device type *type;

    struct bus type *bus; /* type of bus device is on */
    struct device driver *driver; /* which driver has allocated this
                                   device */
    void *platform_data; /* Platform specific data, device
                           core doesn't touch it */
    void *driver_data; /* Driver data, set and get with
                           dev_set_drvdata/dev_get_drvdata */

#ifdef CONFIG_PROVE_LOCKING
    struct mutex lockdep_mutex;
#endif

    struct mutex mutex; /* mutex to synchronize calls to
                           * its driver.
                           */

    struct dev links info links;
    struct dev_pm_info power;
    struct dev_pm domain *pm_domain;

#ifdef CONFIG_ENERGY_MODEL
    struct em_perf domain *em_pd;
#endif

#ifdef CONFIG_PINCTRL
    struct dev_pin_info *pins;
#endif

    struct dev_msi_info msi;
#ifdef CONFIG_DMA_OPS
    const struct dma_map_ops *dma_ops;
#endif

    u64 *dma_mask; /* dma mask (if dma'able device) */
    u64 coherent_dma_mask; /* Like dma_mask, but for
                               alloc_coherent mappings as
                               not all hardware supports
                               64 bit addresses for consistent
                               allocations such descriptors. */

    u64 bus_dma_limit; /* upstream dma constraint */
    const struct bus_dma_region *dma_range_map;

    struct device_dma_parameters *dma_parms;

    struct list_head dma_pools; /* dma pools (if dma'ble) */

#ifdef CONFIG_DMA_DECLARE_COHERENT
    struct dma_coherent_mem *dma_mem; /* internal for coherent mem
                                           override */
#endif

#ifdef CONFIG_DMA_CMA
    struct cma *cma_area; /* contiguous memory area for dma

```

```

allocations */

#endif
#ifdef CONFIG_SWIOTLB
    struct io_tlb_mem *dma_io_tlb_mem;
#endif

/* arch specific additions */
struct dev_archdata archdata;

struct device_node *of_node; /* associated device tree node */
struct fwnode_handle *fwnode; /* firmware device node */

#ifdef CONFIG_NUMA
    int numa_node; /* NUMA node this device is close to */
#endif

    dev_t devt; /* dev_t, creates the sysfs "dev" */
    u32 id; /* device instance */

    spinlock_t devres_lock;
    struct list_head devres_head;

    struct class *class;
    const struct attribute_group *groups; /* optional groups */

    void (*release)(struct device *dev);
    struct iommu_group *iommu_group;
    struct dev_iommu *iommu;

    enum device_removable removable;

    bool offline_disabled:1;
    bool offline:1;
    bool of_node_reused:1;
    bool state_synced:1;
    bool can_match:1;
#ifdef defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_DEVICE) || \
    defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_CPU) || \
    defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_CPU_ALL)
    bool dma_coherent:1;
#endif
#ifdef CONFIG_DMA_OPS_BYPASS
    bool dma_ops_bypass:1;
#endif

};

struct device_driver {
    const char *name;
    struct bus_type *bus;

    struct module *owner;
    const char *mod_name; /* used for built-in modules */

    bool suppress_bind_attrs; /* disables bind/unbind via sysfs */
    enum probe_type probe_type;

    const struct of_device_id *of_match_table;
    const struct acpi_device_id *acpi_match_table;

    int (*probe)(struct device *dev);

```

```

void (*sync_state)(struct device *dev);
int (*remove) (struct device *dev);
void (*shutdown) (struct device *dev);
int (*suspend) (struct device *dev, pm_message_t state);
int (*resume) (struct device *dev);
const struct attribute_group **groups;
const struct attribute_group **dev_groups;

const struct dev_pm_ops *pm;
void (*coredump) (struct device *dev);

struct driver_private *p;
};

#define MAJOR(dev)      ((unsigned int) ((dev) >> MINORBITS))
#define MINOR(dev)     ((unsigned int) ((dev) & MINORMASK))
#define MKDEV(ma,mi)  (((ma) << MINORBITS) | (mi))

* @from: the first in the desired range of device numbers; must include the major number.
* @count: the number of consecutive device numbers required
* @name: the name of the device or driver.
* Return value is zero on success, a negative error code on failure.
int register_chrdev_region(dev_t from, unsigned count, const char *name)

* @dev: output parameter for first assigned number
* @baseminor: first of the requested range of minor numbers
* @count: the number of minor numbers required
* @name: the name of the associated device or driver
int alloc_chrdev_region(dev_t *dev, unsigned baseminor, unsigned count,
                        const char *name)

* @from: the first in the range of numbers to unregister
* @count: the number of device numbers to unregister
*
* This function will unregister a range of @count device numbers,
* starting with @from. The caller should normally be the one who
* allocated those numbers in the first place...
void unregister_chrdev_region(dev_t from, unsigned count)

request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags,
             const char *name, void *dev)
const void *free_irq(unsigned int irq, void *dev_id)

#IRQF_SHARED - allow sharing the irq among several devices
struct tasklet_struct
{
    struct tasklet_struct *next;
    unsigned long state;
    atomic_t count;
    bool use_callback;
    union {
        void (*func)(unsigned long data);
        void (*callback)(struct tasklet_struct *t);
    };
    unsigned long data;
};

```

```

#define DECLARE_TASKLET(name, callback) \
struct tasklet_struct name = { \
    .count = ATOMIC_INIT(0), \
    .callback = callback, \
    .use_callback = true, \
}

enum
{
    TASKLET_STATE_SCHED, /* Tasklet is scheduled for execution */
    TASKLET_STATE_RUN /* Tasklet is running (SMP only) */
};

static inline void tasklet_schedule(struct tasklet_struct *t);
static inline void tasklet_hi_schedule(struct tasklet_struct *t);
static inline void tasklet_disable(struct tasklet_struct *t);
static inline void tasklet_enable(struct tasklet_struct *t);
void tasklet_init(struct tasklet_struct *t, void (*func)(unsigned long), unsigned long data);
void tasklet_kill(struct tasklet_struct *t);

struct block_device {
    sector_t bd_start_sect;
    sector_t bd_nr_sectors;
    struct disk_stats __percpu *bd_stats;
    unsigned long bd_stamp;
    bool bd_read_only; /* read-only policy */
    dev_t bd_dev;
    int bd_openers;
    struct inode * bd_inode; /* will die */
    struct super_block * bd_super;
    void * bd_claiming;
    struct device bd_device;
    void * bd_holder;
    int bd_holders;
    bool bd_write_holder;
    struct kobject * bd_holder_dir;
    u8 bd_partno;
    spinlock_t bd_size_lock; /* for bd_inode->i_size updates */
    struct gendisk * bd_disk;
    struct request_queue * bd_queue;

    /* The counter of freeze processes */
    int bd_fsfreeze_count;
    /* Mutex for freeze */
    struct mutex bd_fsfreeze_mutex;
    struct super_block * bd_fsfreeze_sb;

    struct partition_meta_info * bd_meta_info;
#ifdef CONFIG_FAIL_MAKE_REQUEST
    bool bd_make_it_fail;
#endif
} randomize_layout;

struct cdev {
    struct kobject kobj;
    struct module * owner;
    const struct file_operations * ops;
    struct list_head list;

```

```

    dev_t dev;
    unsigned int count;
} randomize_layout;

struct inode {
    umode_t                i_mode; // mode
    unsigned short        i_opflags;
    kuid_t                i_uid; // uid
    kgid_t                i_gid; // gid
    unsigned int           i_flags;

#ifdef CONFIG_FS_POSIX_ACL
    struct posix_acl      *i_acl;
    struct posix_acl      *i_default_acl;
#endif

    const struct inode_operations *i_op;
    struct super_block *i_sb; // sb
    struct address_space *i_mapping;

#ifdef CONFIG_SECURITY
    void *i_security;
#endif

    /* Stat data, not accessed from path walking */
    unsigned long i_ino;
    /*
     * Filesystems may only read i_nlink directly. They shall use the
     * following functions for modification:
     *
     * (set|clear|inc|drop)_nlink
     * inode_(inc|dec)_link_count
     */
    union {
        const unsigned int i_nlink;
        unsigned int __i_nlink;
    };
    dev_t                i_rdev;
    loff_t               i_size;
    struct timespec64 i_atime;
    struct timespec64 i_mtime;
    struct timespec64 i_ctime;
    spinlock_t          i_lock; /* i_blocks, i_bytes, maybe i_size */
    unsigned short      i_bytes;
    u8                 i_blkbits;
    u8                 i_write_hint;
    blkcnt_t           i_blocks;

#ifdef __NEED_I_SIZE_ORDERED
    seqcount_t         i_size_seqcount;
#endif

    /* Misc */
    unsigned long i_state;
    struct rw_semaphore i_rwsem;

    unsigned long dirtied_when; /* jiffies of first dirtying */

```

```

unsigned long                dirtied_time_when;

struct hlist_node    i_hash;
struct list_head     i_io_list; /* backing dev IO list */
#ifdef CONFIG_CGROUP_WRITEBACK
struct bdi_writeback    *i_wb;          /* the associated cgroup wb */

/* foreign inode detection, see wbc_detach_inode() */
int                    i_wb_frn_winner;
u16                   i_wb_frn_avg_time;
u16                   i_wb_frn_history;
#endif

struct list_head      i_lru;            /* inode LRU list */
struct list_head      i_sb_list;
struct list_head      i_wb_list; /* backing dev writeback list */
union {
    struct hlist_head    i_dentry;
    struct rcu_head      i_rcu;
};
atomic64_t            i_version;
atomic64_t            i_sequence; /* see futex */
atomic_t              i_count;
atomic_t              i_dio_count;
atomic_t              i_writecount;
#ifdef defined(CONFIG_IMA) || defined(CONFIG_FILE_LOCKING)
atomic_t              i_readcount; /* struct files open RO */
#endif

union {
    const struct file_operations    *i_fop; /* former ->i_op->default_file_ops */
    void (*free_inode)(struct inode *);
};
struct file_lock_context    *i_flctx;
struct address_space        i_data;
struct list_head            i_devices;
union {
    struct pipe_inode_info    *i_pipe;
    struct cdev               *i_cdev;
    char                     *i_link;
    unsigned                 i_dir_seq;
};

__u32                    i_generation;

#ifdef CONFIG_FSNOTIFY
__u32                    i_fsnotify_mask; /* all events this inode cares about */
struct fsnotify_mark_connector __rcu    *i_fsnotify_marks;
#endif

#ifdef CONFIG_FS_ENCRYPTION
struct fscrypt_info    *i_crypt_info;
#endif

#ifdef CONFIG_FS_VERITY
struct fsverity_info    *i_verity_info;
#endif

void                    *i_private; /* fs or device private pointer */

```

```

} __randomize_layout;

struct workqueue_struct {
    struct list_head    pwqs;          /* WR: all pwqs of this wq */
    struct list_head    list;          /* PR: list of all workqueues */

    struct mutex          mutex;          /* protects this wq */
    int                  work_color;    /* WQ: current work color */
    int                  flush_color;    /* WQ: current flush color */
    atomic_t             nr_pwqs_to_flush; /* flush in progress */
    struct wq_flusher    *first_flusher; /* WQ: first flusher */
    struct list_head    flusher_queue; /* WQ: flush waiters */
    struct list_head    flusher_overflow; /* WQ: flush overflow list */

    struct list_head    maydays; /* MD: pwqs requesting rescue */
    struct worker       *rescuer; /* MD: rescue worker */

    int                  nr_drainers;    /* WQ: drain in progress */
    int                  saved_max_active; /* WQ: saved pwq max_active */

    struct workqueue_attrs *unbound_attrs; /* PW: only for unbound wqs */
    struct pool_workqueue *dfl_pwq; /* PW: only for unbound wqs */

#ifdef CONFIG_SYSFS
    struct wq_device    *wq_dev; /* I: for sysfs interface */
#endif
#ifdef CONFIG_LOCKDEP
    char                  *lock_name;
    struct lock_class_key key;
    struct lockdep_map lockdep_map;
#endif

    char                  name[WQ_NAME_LEN]; /* I: workqueue name */

    /*
     * Destruction of workqueue_struct is RCU protected to allow walking
     * the workqueues list without grabbing wq_pool_mutex.
     * This is used to dump all workqueues from sysrq.
     */
    struct rcu_head      rcu;

    /* hot fields used during command issue, aligned to cacheline */
    unsigned int          flags cacheline_aligned; /* WQ: WQ_* flags */
    struct pool_workqueue percpu *cpu_pwqs; /* I: per-cpu pwqs */
    struct pool_workqueue rcu *numa_pwq_tbl[]; /* PWR: unbound pwqs indexed by node */
};

struct work_struct {
    atomic_long_t data;
    struct list_head entry;
    work_func_t func;
#ifdef CONFIG_LOCKDEP
    struct lockdep_map lockdep_map;
#endif
};

#define DECLARE_WORK(name, void (*f)(void *))
#define INIT_WORK(struct work_struct *work, void (*f)(void), void *)
static inline bool queue_work(struct workqueue_struct *wq, struct work_struct *work);

```

```

void flush_workqueue(struct workqueue_struct *wq);
extern void destroy_workqueue(struct workqueue_struct *wq);
void open_softirq(int nr, void (*action)(struct softirq_action *))
{
    softirq_vec[nr].action = action;
}
void raise_softirq(unsigned int nr)
{
    unsigned long flags;

    local_irq_save(flags);
    raise_softirq_irqoff(nr);
    local_irq_restore(flags);
}
struct proc_dir_entry {
    /*
     * number of callers into module in progress;
     * negative -> it's going away RSN
     */
    atomic_t in_use;
    refcount_t refcnt;
    struct list_head pde_openers; /* who did ->open, but not ->release */
    /* protects ->pde_openers and all struct pde_opener instances */
    spinlock_t pde_unload_lock;
    struct completion *pde_unload_completion;
    const struct inode_operations *proc_iops;
    union {
        const struct proc_ops *proc_ops;
        const struct file_operations *proc_dir_ops;
    };
    const struct dentry_operations *proc_dops;
    union {
        const struct seq_operations *seq_ops;
        int (*single_show)(struct seq_file *, void *);
    };
    proc_write_t write;
    void *data;
    unsigned int state_size;
    unsigned int low_ino;
    nlink_t nlink;
    kuid_t uid;
    kgid_t gid;
    loff_t size;
    struct proc_dir_entry *parent;
    struct rb_root subdir;
    struct rb_node subdir_node;
    char *name;
    umode_t mode;
    u8 flags;
    u8 namelen;
    char inline_name[];
}
randomize_layout;
struct proc_dir_entry *proc_create(const char *name, umode_t mode, struct proc_dir_entry *parent,
const struct proc_ops *proc_ops);
extern struct proc_dir_entry *proc_mkdir(const char *, struct proc_dir_entry *);
extern struct proc_dir_entry *proc_symlink(const char *,
struct proc_dir_entry *, const char *);

```



```

extern void remove_proc_entry(const char *, struct proc_dir_entry *);
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
    ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
    int (*iopoll)(struct kiocb *kiocb, struct io_comp_batch *,
                  unsigned int flags);
    int (*iterate) (struct file *, struct dir_context *);
    int (*iterate_shared) (struct file *, struct dir_context *);
    __poll_t (*poll) (struct file *, struct poll_table_struct *);
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    unsigned long mmap_supported_flags;
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, loff_t, loff_t, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
    unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned long, unsigned
long, unsigned long);
    int (*check_flags)(int);
    int (*flock) (struct file *, int, struct file_lock *);
    ssize_t (*splice_write)(struct pipe_inode_info *, struct file *, loff_t *, size_t, unsigned int);
    ssize_t (*splice_read)(struct file *, loff_t *, struct pipe_inode_info *, size_t, unsigned int);
    int (*setlease)(struct file *, long, struct file_lock **, void **);
    long (*fallocate)(struct file *file, int mode, loff_t offset,
                      loff_t len);
    void (*show_fdinfo)(struct seq_file *m, struct file *f);
#ifdef CONFIG_MMU
    unsigned (*mmap_capabilities)(struct file *);
#endif
    ssize_t (*copy_file_range)(struct file *, loff_t, struct file *,
                              loff_t, size_t, unsigned int);
    loff_t (*remap_file_range)(struct file *file_in, loff_t pos_in,
                              struct file *file_out, loff_t pos_out,
                              loff_t len, unsigned int remap_flags);
    int (*fadvise)(struct file *, loff_t, loff_t, int);
} __randomize_layout;
unsigned long copy_to_user(void __user *to, const void *from, unsigned long n);

struct file_system_type {
    const char *name; /* = name */
    int fs_flags;
#define FS_REQUIRES_DEV 1
#define FS_BINARY_MOUNTDATA 2
#define FS_HAS_SUBTYPE 4
#define FS_USERNS_MOUNT 8 /* Can be mounted by userns root */
#define FS_DISALLOW_NOTIFY_PERM 16 /* Disable fanotify permission events */
#define FS_ALLOW_IDMAP 32 /* FS has been updated to handle vfs idmappings. */
#define FS_RENAME_DOES_D_MOVE 32768 /* FS will handle d_move() during
rename() internally. */

```

```

int (*init fs context)(struct fs context *);
const struct fs parameter spec *parameters;
struct dentry *(mount) (struct file system type *, int,
                        const char *, void *); /* = mount */
void (*kill_sb) (struct super_block *); /* = killsb */
struct module *owner; /* = this module */
struct file system type * next;
struct hlist head fs supers;

struct lock class key s lock key;
struct lock class key s umount key;
struct lock class key s vfs rename key;
struct lock class key s writers key[SB_FREEZE_LEVELS];

struct lock class key i lock key;
struct lock class key i mutex key;
struct lock class key invalidate lock key;
struct lock class key i mutex dir key;
};

struct super_block {
    struct list_head      s_list;           /* Keep this first */
    dev_t                  s_dev;           /* search index; _not_kdev_t */
    unsigned char          s_blocksize_bits; /* = page_shift */
    unsigned long          s_blocksize; /* = page_size */
    loff_t                 s_maxbytes;      /* Max file size */
    struct file_system_type *s_type; // type
    const struct super_operations *s_op; /* = ops */
    const struct dquot_operations *dq_op;
    const struct quotactl_ops *s_qcop;
    const struct export_operations *s_export_op;
    unsigned long          s_flags;
    unsigned long          s_iflags; /* internal SB_I_* flags */
    unsigned long          s_magic; /* = magic */
    struct dentry          *s_root; /* = root */
    struct rw_semaphore    s_umount;
    int                    s_count;
    atomic_t                s_active;
#ifdef CONFIG_SECURITY
    void                    *s_security;
#endif
    const struct xattr_handler **s_xattr;
#ifdef CONFIG_FS_ENCRYPTION
    const struct fscrypt_operations *s_cop;
    struct key *s_master_keys; /* master crypto keys in use */
#endif
#ifdef CONFIG_FS_VERITY
    const struct fsverity_operations *s_vop;
#endif
#ifdef IS_ENABLED(CONFIG_UNICODE)
    struct unicode_map *s_encoding;
    __u16 s_encoding_flags;
#endif
    struct hlist_bl_head s_roots; /* alternate root dentries for NFS */
    struct list_head      s_mounts; /* list of mounts; _not_for fs use */
    struct block_device *s_bdev;
    struct backing_dev_info *s_bdi;

```

```

struct mtd_info          *s_mtd;
struct hlist_node        s_instances;
unsigned int             s_quota_types;    /* Bitmask of supported quota types */
struct quota_info        s_dquot; /* Diskquota specific options */

struct sb_writers        s_writers;

/*
 * Keep s_fs_info, s_time_gran, s_fsnotify_mask, and
 * s_fsnotify_marks together for cache efficiency. They are frequently
 * accessed and rarely modified.
 */
void                    *s_fs_info;        /* Filesystem private info */

/* Granularity of c/m/atime in ns (cannot be worse than a second) */
u32                     s_time_gran;
/* Time limits for c/m/atime in seconds */
time64_t                s_time_min;
time64_t                s_time_max;
#ifdef CONFIG_FSNOTIFY
__u32                   s_fsnotify_mask;
struct fsnotify_mark_connector __rcu *s_fsnotify_marks;
#endif

char                    s_id[32]; /* Informational name */
uuid_t                  s_uuid;    /* UUID */

unsigned int            s_max_links;
fmode_t                 s_mode;

/*
 * The next field is for VFS *only*. No filesystems have any business
 * even looking at it. You had been warned.
 */
struct mutex s_vfs_rename_mutex;    /* Kludge */

/*
 * Filesystem subtype. If non-empty the filesystem type field
 * in /proc/mounts will be "type.subtype"
 */
const char *s_subtype;

const struct dentry_operations *s_d_op; /* default d_op for dentries */

struct shrinker s_shrink;    /* per-sb shrinker handle */

/* Number of inodes with nlink == 0 but still referenced */
atomic_long_t s_remove_count;

/*
 * Number of inode/mount/sb objects that are being watched, note that
 * inodes objects are currently double-accounted.
 */
atomic_long_t s_fsnotify_connectors;

/* Being remounted read-only */
int s_readonly_remount;

```

```

/* per-sb errseq_t for reporting writeback errors via syncfs */
errseq_t s_wb_err;

/* AIO completions deferred from interrupt context */
struct workqueue_struct *s_dio_done_wq;
struct hlist_head s_pins;

/*
 * Owning user namespace and default context in which to
 * interpret filesystem uids, gids, quotas, device nodes,
 * xattrs and security labels.
 */
struct user_namespace *s_user_ns;

/*
 * The list_lru structure is essentially just a pointer to a table
 * of per-node lru lists, each of which has its own spinlock.
 * There is no need to put them into separate cachelines.
 */
struct list_lru          s_dentry_lru;
struct list_lru          s_inode_lru;
struct rcu_head          rcu;
struct work_struct destroy_work;

struct mutex              s_sync_lock;          /* sync serialisation lock */

/*
 * Indicates how deep in a filesystem stack this SB is
 */
int s_stack_depth;

/* s_inode_list_lock protects s_inodes */
spinlock_t               s_inode_list_lock ____cacheline_aligned_in_smp;
struct list_head          s_inodes; /* all inodes */

spinlock_t               s_inode_wblist_lock;
struct list_head          s_inodes_wb; /* writeback inodes */
} __randomize_layout;

struct dentry {
/* RCU lookup touched fields */
unsigned int d_flags; /* protected by d_lock */
seqcount_spinlock_t d_seq; /* per dentry seqlock */
struct hlist_bl_node d_hash; /* lookup hash list */
struct dentry *d_parent; /* parent directory */
struct qstr d_name;
struct inode *d_inode; /* Where the name belongs to - NULL is
                        * negative */
unsigned char d_iname[DNAME_INLINE_LEN]; /* small names */

/* Ref lookup also touches following */
struct lockref d_lockref; /* per-dentry lock and refcount */
const struct dentry_operations *d_op;
struct super_block *d_sb; /* The root of the dentry tree */
unsigned long d_time; /* used by d_revalidate */
void *d_fsdata; /* fs-specific data */

```

```

union {
    struct list head d_lru;          /* LRU list */
    struct wait queue head t *d_wait; /* in-lookup ones only */
};
struct list head d_child;          /* child of parent list */
struct list head d_subdirs;        /* our children */
/*
 * d_alias and d_rcu can share memory
 */
union {
    struct hlist node d_alias;      /* inode alias list */
    struct hlist bl node d_in_lookup_hash; /* only for in-lookup ones */
    struct rcu head d_rcu;
} d_u;
} randomize layout;

struct file {
    union {
        struct list node fu_llist;
        struct rcu head fu_rcuhead;
    } f_u;
    struct path f_path;             /* path
    struct inode *f_inode;          /* cached value: inode */
    const struct file_operations *f_op;

    /*
     * Protects f_ep, f_flags.
     * Must not be taken from IRQ context.
     */
    struct spinlock t f_lock;
    struct atomic long t f_count;
    unsigned int f_flags;
    struct fmode t f_mode;          /* mode
    struct mutex f_pos_lock;
    struct loff t f_pos;           /* pos
    struct fown_struct f_owner;
    const struct cred *f_cred;
    struct file_ra_state f_ra;

    struct u64 f_version;

#ifdef CONFIG_SECURITY
    void *f_security;
#endif

    /* needed for tty driver, and maybe others */
    void *private_data;

#ifdef CONFIG_EPOLL
    /* Used by fs/eventpoll.c to link all the hooks to this file */
    struct hlist head *f_ep;
#endif /* #ifdef CONFIG_EPOLL */
    struct address_space *f_mapping;
    struct errseq t f_wb_err;
    struct errseq t f_sb_err; /* for syncfs */
} randomize layout
attribute ((aligned(4))); /* lest something weird decides that 2 is OK */

```

```

struct file handle {
    u32 handle bytes;
    int handle type;
    /* file identifier */
    unsigned char f handle[];
};

extern struct dentry *mount bdev(struct file system type *fs type,
    int flags, const char *dev name, void *data,
    int (*fill super)(struct super block *, void *, int));
extern struct dentry *mount nodev(struct file system type *fs type,
    int flags, void *data,
    int (*fill super)(struct super block *, void *, int));
struct kmem cache *kmem cache create(const char *name, unsigned int size,
    unsigned int align, slab flags t flags,
    void (*ctor)(void *));
void kmem cache destroy(struct kmem cache *s);
void *kmem cache alloc(struct kmem cache *s, gfp t flags) // GFP_KERNEL
void kmem cache free(struct kmem cache *s, void *objp);
extern int register filesystem(struct file system type *);
extern int unregister filesystem(struct file system type *);
extern struct inode *new inode(struct super block *sb);
struct super operations {
    struct inode *(*alloc inode)(struct super block *sb);
    void (*destroy inode)(struct inode *);
    void (*free inode)(struct inode *);

    void (*dirty inode) (struct inode *, int flags);
    int (*write inode) (struct inode *, struct writeback control *wbc);
    int (*drop inode) (struct inode *) = generic_delete_inode;
    void (*evict inode) (struct inode *);
    void (*put super) (struct super block *);
    int (*sync fs)(struct super block *sb, int wait);
    int (*freeze super) (struct super block *);
    int (*freeze fs) (struct super block *);
    int (*thaw super) (struct super block *);
    int (*unfreeze fs) (struct super block *);
    int (*statfs) (struct dentry *, struct kstatfs *) = simple_stat_fs;
    int (*remount fs) (struct super block *, int *, char *);
    void (*umount begin) (struct super block *);

    int (*show options)(struct seq file *, struct dentry *);
    int (*show devname)(struct seq file *, struct dentry *);
    int (*show path)(struct seq file *, struct dentry *);
    int (*show stats)(struct seq file *, struct dentry *);
#ifdef CONFIG QUOTA
    ssize t (*quota read)(struct super block *, int, char *, size t, loff t);
    ssize t (*quota write)(struct super block *, int, const char *, size t, loff t);
    struct dquot **(*get dquots)(struct inode *);
#endif
    long (*nr cached objects)(struct super block *,
        struct shrink control *);
    long (*free cached objects)(struct super block *,
        struct shrink control *);
};
int generic delete inode(struct inode *inode)
{
    return 1;
}

```

```

}
int simple_statfs(struct dentry *dentry, struct kstatfs *buf)
{
    buf->f_type = dentry->d_sb->s_magic;
    buf->f_bsize = PAGE_SIZE;
    buf->f_namelen = NAME_MAX;
    return 0;
}

struct inode_operations {
    struct dentry * (*lookup) (struct inode *, struct dentry *, unsigned int);
    const char * (*get_link) (struct dentry *, struct inode *, struct delayed_call *);
    int (*permission) (struct user_namespace *, struct inode *, int);
    struct posix_acl * (*get_acl) (struct inode *, int, bool);

    int (*readlink) (struct dentry *, char __user *, int);

    int (*create) (struct user_namespace *, struct inode *, struct dentry *,
                  u_mode_t, bool);
    int (*link) (struct dentry *, struct inode *, struct dentry *);
    int (*unlink) (struct inode *, struct dentry *);
    int (*symlink) (struct user_namespace *, struct inode *, struct dentry *,
                  const char *);
    int (*mkdir) (struct user_namespace *, struct inode *, struct dentry *,
                  u_mode_t);
    int (*rmdir) (struct inode *, struct dentry *);
    int (*mknod) (struct user_namespace *, struct inode *, struct dentry *,
                  u_mode_t, dev_t);
    int (*rename) (struct user_namespace *, struct inode *, struct dentry *,
                  struct inode *, struct dentry *, unsigned int);
    int (*setattr) (struct user_namespace *, struct dentry *,
                  struct iattr *);
    int (*getattr) (struct user_namespace *, const struct path *,
                  struct kstat *, u32, unsigned int);
    ssize_t (*listxattr) (struct dentry *, char *, size_t);
    int (*fiemap) (struct inode *, struct fiemap_extent_info *, u64 start,
                  u64 len);
    int (*update_time) (struct inode *, struct timespec64 *, int);
    int (*atomic_open) (struct inode *, struct dentry *,
                        struct file *, unsigned open_flag,
                        u_mode_t create_mode);
    int (*tmpfile) (struct user_namespace *, struct inode *,
                  struct dentry *, u_mode_t);
    int (*set_acl) (struct user_namespace *, struct inode *,
                  struct posix_acl *, int);
    int (*fileattr_set) (struct user_namespace * mnt_userns,
                        struct dentry * dentry, struct fileattr * fa);
    int (*fileattr_get) (struct dentry * dentry, struct fileattr * fa);
} ____cacheline_aligned;

struct dentry_operations {
    int (*d_revalidate) (struct dentry *, unsigned int);
    int (*d_weak_revalidate) (struct dentry *, unsigned int);
    int (*d_hash) (const struct dentry *, struct qstr *);
    int (*d_compare) (const struct dentry *,
                      unsigned int, const char *, const struct qstr *);
    int (*d_delete) (const struct dentry *);
    int (*d_init) (struct dentry *);
    void (*d_release) (struct dentry *);

```

```

void (*d_prune)(struct dentry *);
void (*d_iput)(struct dentry *, struct inode *);
char *(*d_dname)(struct dentry *, char *, int);
struct vfsmount *(*d_automount)(struct path *);
int (*d_manage)(const struct path *, bool);
struct dentry *(*d_real)(struct dentry *, const struct inode *);
} cacheline_aligned;

struct socket {
    socket_state                state;

    short                        type;

    unsigned long                flags;

    struct file                  *file;
    struct sock                  *sk;
    const struct proto_ops       *ops;

    struct socket_wq            wq;
};

typedef enum {
    SS_FREE = 0,                /* not allocated */
    SS_UNCONNECTED,            /* unconnected to any socket */
    SS_CONNECTING,            /* in process of connecting */
    SS_CONNECTED,            /* connected to socket */
    SS_DISCONNECTING         /* in process of disconnecting */
} socket_state;

struct sockaddr {
    sa_family_t                sa_family;    /* address family, AF_XXX */
    char                        sa_data[14];    /* 14 bytes of protocol address */
};

struct sockaddr_in {
    kernel sa_family_t        sin_family;    /* Address family */
    be16                      sin_port;        /* Port number */
    struct in_addr            sin_addr;        /* Internet address */

    /* Pad to size of `struct sockaddr'. */
    unsigned char                pad[SOCK_SIZE - sizeof(short int) -
                                     sizeof(unsigned short int) - sizeof(struct in_addr)];
};

Sys_socket -> sock_create
#define UNIX_PATH_MAX 108
struct sockaddr_un {
    kernel sa_family_t sun_family; /* AF_UNIX */
    char sun_path[UNIX_PATH_MAX]; /* pathname */
};

#define htonl(x) cpu to be32(x)
#define htons(x) cpu to be16(x)
#define ntohl(x) be32 to cpu(x)
#define ntohs(x) be16 to cpu(x)

#define htonl(x) htonl(x)
#define ntohl(x) ntohl(x)
#define htons(x) htons(x)
#define ntohs(x) ntohs(x)

```



```

struct proto_ops {
    int                family;
    struct module      *owner;
    int                (*release) (struct socket *sock);
    int                (*bind)    (struct socket *sock,
                                   struct sockaddr *myaddr,
                                   int sockaddr_len);
    int                (*connect) (struct socket *sock,
                                   struct sockaddr *vaddr,
                                   int sockaddr_len, int flags);
    int                (*socketpair)(struct socket *sock1,
                                      struct socket *sock2);
    int                (*accept)  (struct socket *sock,
                                   struct socket *newssock, int flags, bool kern);
    int                (*getname) (struct socket *sock,
                                   struct sockaddr *addr,
                                   int peer);
    __poll_t          (*poll)    (struct file *file, struct socket *sock,
                                   struct poll_table_struct *wait);
    int                (*ioctl)  (struct socket *sock, unsigned int cmd,
                                   unsigned long arg);
#ifdef CONFIG_COMPAT
    int                (*compat_ioctl) (struct socket *sock, unsigned int cmd,
                                         unsigned long arg);
#endif
    int                (*gettstamp) (struct socket *sock, void __user *userstamp,
                                      bool timeval, bool time32);
    int                (*listen)  (struct socket *sock, int len);
    int                (*shutdown) (struct socket *sock, int flags);
    int                (*setsockopt)(struct socket *sock, int level,
                                      int optname, sockptr_t optval,
                                      unsigned int optlen);
    int                (*getsockopt)(struct socket *sock, int level,
                                      int optname, char __user *optval, int __user *optlen);
    void                (*show_fdinfo)(struct seq_file *m, struct socket *sock);
    int                (*sendmsg) (struct socket *sock, struct msghdr *m,
                                   size_t total_len);
    /* Notes for implementing recvmsg:
     * =====
     * msg->msg_namelen should get updated by the recvmsg handlers
     * iff msg_name != NULL. It is by default 0 to prevent
     * returning uninitialized memory to user space. The recvfrom
     * handlers can assume that msg.msg_name is either NULL or has
     * a minimum size of sizeof(struct sockaddr_storage).
     */
    int                (*recvmsg) (struct socket *sock, struct msghdr *m,
                                   size_t total_len, int flags);
    int                (*mmap)    (struct file *file, struct socket *sock,
                                   struct vm_area_struct *vma);
    ssize_t            (*sendpage) (struct socket *sock, struct page *page,
                                    int offset, size_t size, int flags);
    ssize_t            (*splice_read)(struct socket *sock, loff_t *ppos,
                                       struct pipe_inode_info *pipe, size_t len, unsigned int
flags);
    int                (*set_peek_off)(struct sock *sk, int val);
    int                (*peek_len)(struct socket *sock);

```

```

/* The following functions are called internally by kernel with
 * sock lock already held.
 */
int (*read_sock)(struct sock *sk, read_descriptor_t *desc,
                 sk_read_actor_t recv_actor);
int (*sendpage_locked)(struct sock *sk, struct page *page,
                       int offset, size_t size, int flags);
int (*sendmsg_locked)(struct sock *sk, struct msghdr *msg,
                      size_t size);
int (*set_rcvlowat)(struct sock *sk, int val);
};

extern struct task_struct init_task;
#define next_task(p) \
    list_entry_rcu((p)->tasks.next, struct task_struct, tasks)

struct task_struct {
    /* describes a process running in the system, created dynamically. */
    int prio;
    int static_prio;
    char comm[TASK_COMM_LEN]; /* executable name excluding path
    struct list_head tasks;
    ...
    struct mm_struct *mm;
    struct mm_struct *active_mm;
    ...
    pid_t pid;
    pid_t tgid;
    ...
    /*
    * Children/sibling form the list of natural children:
    */
    struct list_head children;
    struct list_head sibling;
    ...
    /* Filesystem information: */
    struct fs_struct *fs;

    /* Open file information: */
    struct files_struct *files;
    /* Namespaces: */
    struct nsproxy *nsproxy;
};

struct fs_struct {
    /* information about the file system to which the process belongs */
    int users;
    spinlock_t lock;
    seqcount spinlock_t seq;
    int umask;
    int in_exec;
    struct path root, pwd; the mounting object of the root directory and working directory
} __randomize_layout;
/*
 * Open file table structure

```

```

*/
struct files_struct {
/*
    * read mostly part
    */
    atomic_t count;
    bool resize_in_progress;
    wait_queue_head_t resize_wait;

    struct fdtable __rcu * fdt;
    struct fdtable fdtab;
/*
    * written part on a separate cache line in SMP
    */
    spinlock_t file_lock cacheline_aligned_in_smp;
    unsigned int next_fd;
    unsigned long close_on_exec_init[1]; // fd that should be closed when exec() is called
    unsigned long open_fds_init[1]; // initial set of fd
    unsigned long full_fds_bits_init[1];
    struct file __rcu * fd_array[NR_OPEN_DEFAULT];
};

```